


Concurrency: Deadlock and Starvation

**S.Rajarajan APII/CSE
SASTRA**



Deadlock



- Permanent blocking of a set of processes that either compete for system resources or communicate with each other
- No efficient solution
- Involve conflicting needs for resources by two or more processes

Deadlock




Figure 6.1 Illustration of Deadlock

Deadlock



Process P

• • •

Get A

• • •

Get B

• • •

Release A

• • •

Release B

• • •

Process Q

• • •

Get B

• • •

Get A

• • •

Release B

• • •

Release A

• • •

Deadlock - Joint Progress Diagram

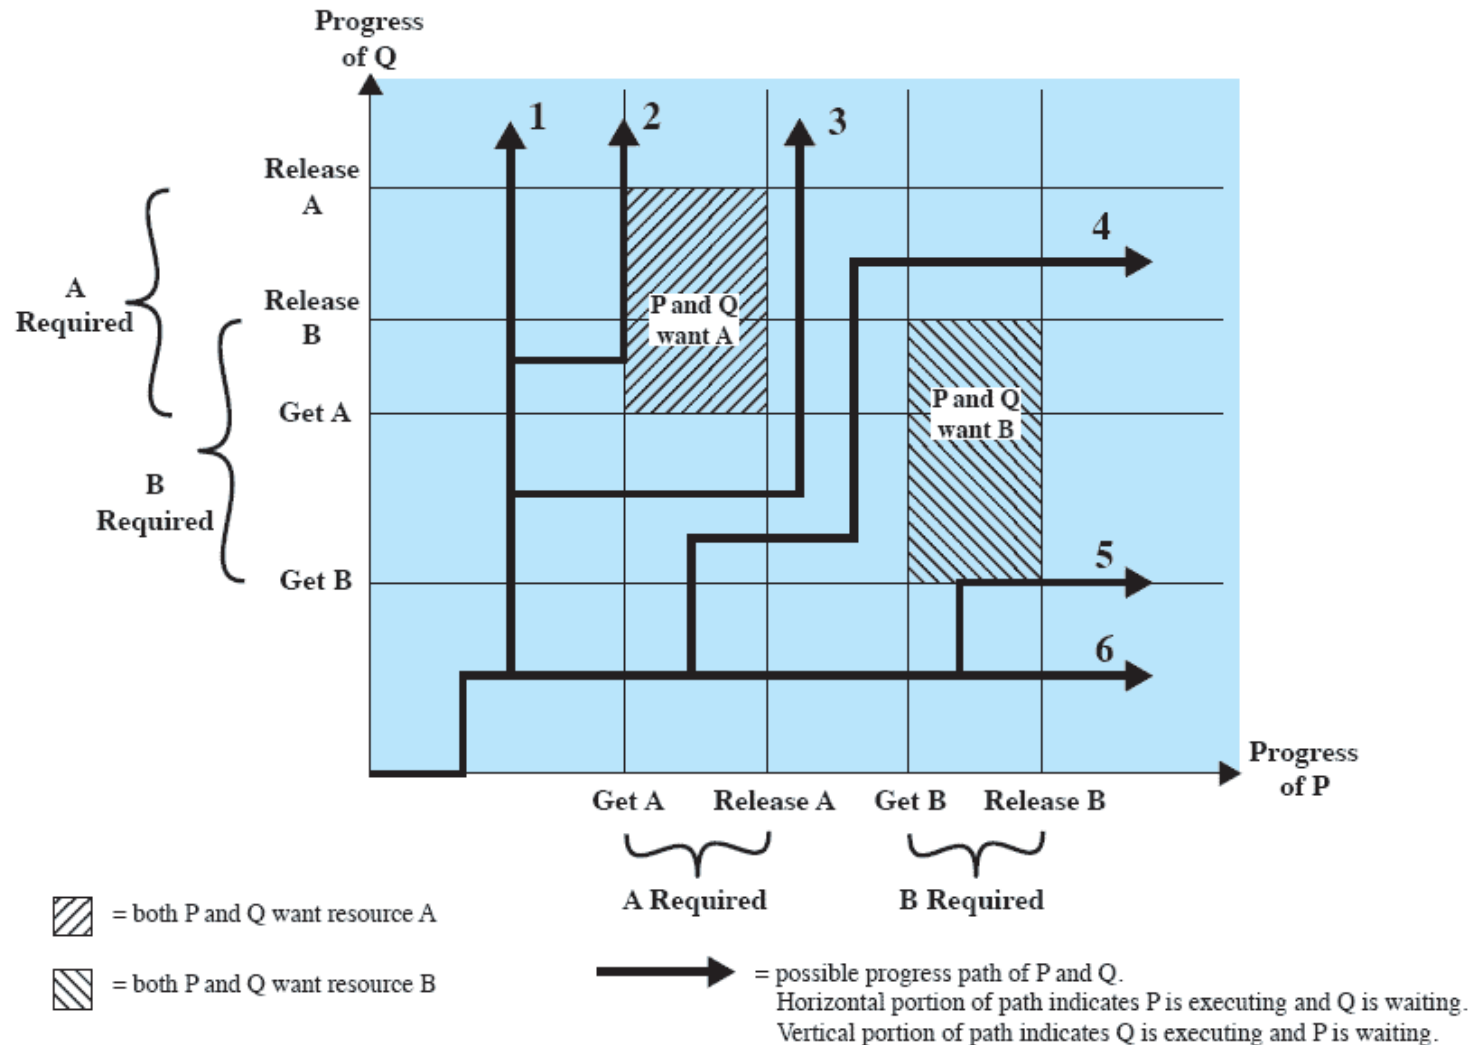


Figure 6.3 Example of No Deadlock [BACO03]

Possible execution Paths



- 1. Q acquires B and then A and then releases B and A.
 - 2. Q acquires B and then A. P executes and blocks for A.
 - 3. Q acquires B and P acquires A. -DL
 - 4. P acquires A and then acquires B, latter releases B and A
 - 5. P acquires A and then acquires B, Q requests B and get blocked
 - 6. P acquires, Q acquires B – DL
-
- The gray shaded area in an ***Joint Progress Diagram*** is called ***Fata Region***

Reusable Resources




- Used by only one process at a time and not depleted by that use
- Processes obtain resources that they later release for reuse by other processes

Reusable Resources



- Processors, I/O channels, main and secondary memory, devices, and data structures such as files, databases, and semaphores
- Deadlock occurs if each process holds one resource and requests the other resource

Reusable Resources



Process P

| Step | Action |
|----------------|------------------|
| p ₀ | Request (D) |
| p ₁ | Lock (D) |
| p ₂ | Request (T) |
| p ₃ | Lock (T) |
| p ₄ | Perform function |
| p ₅ | Unlock (D) |
| p ₆ | Unlock (T) |

Process Q


| Step | Action |
|----------------|------------------|
| q ₀ | Request (T) |
| q ₁ | Lock (T) |
| q ₂ | Request (D) |
| q ₃ | Lock (D) |
| q ₄ | Perform function |
| q ₅ | Unlock (T) |
| q ₆ | Unlock (D) |

Figure 6.4 Example of Two Processes Competing for Reusable Resources

Reusable Resources



- Space is available for allocation of 200Kbytes, and the following sequence of events occur



- Deadlock occurs if both processes progress to their second request

Consumable Resources





- Created (produced) and destroyed (consumed)
- Interrupts, signals, messages, and information in I/O buffers
- Deadlock may occur if a Receive message is blocking
- May take a rare combination of events to cause deadlock

Example of Deadlock



- Deadlock occurs if receives blocking



Resource Allocation Graphs



- Directed graph that depicts a state of the system of resources and processes





(a) Resource is requested




(b) Resource is held

Resource Allocation Graphs



(c) Circular wait



(d) No deadlock

Resource Allocation Graphs




Figure 6.6 Resource Allocation Graph for Figure 6.1b

Conditions for Deadlock



- Mutual exclusion
 - Only one process may use a resource at a time
- Hold-and-wait
 - A process may hold allocated resources while awaiting assignment of others

Conditions for Deadlock



- No preemption
 - No resource can be forcibly removed from a process holding it
- Circular wait
 - A closed chain of processes exists, such that each process holds at least one resource needed by the next process in the chain

Possibility of Deadlock



- Mutual Exclusion
- No preemption
- Hold and wait

Existence of Deadlock



- Mutual Exclusion
- No preemption
- Hold and wait
- Circular wait

Solution for deadlocks



- Prevention
- Avoidance
- Detection

Deadlock Prevention



- Mutual Exclusion
 - Must be supported by the OS
- Hold and Wait
 - Require a process request all of its required resources at one time
 - Inefficient because
 1. A process waiting for all its resources when it could have progressed with available resources
 2. The resources allotted to a process may remain unused for a long time

Deadlock Prevention



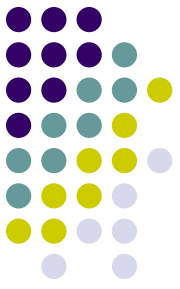
- No Preemption
 - Process must release resource and request again
 - OS may preempt a process to require it releases its resources



- Circular Wait

- Define a linear ordering of resource types
- Assign an index with each resource
- Then Resource R_i precedes R_j in the ordering if $i < j$
- Then if process P_1 holding R_i and demanding R_j while P_2 holding R_j demanding R_i is impossible
- Because, P_2 holding R_j should not ask for R_i
- It should either give up R_j before asking R_i or should ask for some other resource R_{j+1}

Deadlock Avoidance



- Allows the three necessary conditions but makes judicious choices to assure that deadlock point is never reached.
- A decision is made dynamically whether the current resource allocation request will, if granted, potentially lead to a deadlock
- Requires knowledge of future process resource requests

Two Approaches to Deadlock Avoidance



- Do not start a process if its demands might lead to deadlock
- Do not grant an incremental resource request to a process if this allocation might lead to deadlock

Resource Allocation Denial



- Referred to as the banker's algorithm
- State of the system is the current allocation of resources to process
- **Safe state** is where there is at least one sequence that does not result in deadlock
- **Unsafe state** is a state that is not safe


Necessary data structures



| | |
|---|--|
| Resource = $\mathbf{R} = (R_1, R_2, \dots, R_m)$ | Total amount of each resource in the system |
| Available = $\mathbf{V} = (V_1, V_2, \dots, V_m)$ | Total amount of each resource not allocated to any process |
| Claim = $\mathbf{C} = \begin{pmatrix} C_{11} & C_{12} & \dots & C_{1m} \\ C_{21} & C_{22} & \dots & C_{2m} \\ \vdots & \vdots & \ddots & \vdots \\ C_{n1} & C_{n2} & \dots & C_{nm} \end{pmatrix}$ | C_{ij} = requirement of process i for resource j |
| Allocation = $\mathbf{A} = \begin{pmatrix} A_{11} & A_{12} & \dots & A_{1m} \\ A_{21} & A_{22} & \dots & A_{2m} \\ \vdots & \vdots & \ddots & \vdots \\ A_{n1} & A_{n2} & \dots & A_{nm} \end{pmatrix}$ | A_{ij} = current allocation to process i of resource j |

1. $R_j = V_j + \sum_{i=1}^n A_{ij}$, for all j All resources are either available or allocated.
2. $C_{ij} \leq R_j$, for all i, j No process can claim more than the total amount of resources in the system.
3. $A_{ij} \leq C_{ij}$, for all i, j No process is allocated more resources of any type than the process originally claimed to need.

Determination of a Safe State



Start a new process P_{n+1} only if

$$R_j \geq C_{(n+1)j} + \sum_{i=1}^n C_{ij} \quad \text{for all } j$$

| | R1 | R2 | R3 |
|----|----|----|----|
| P1 | 3 | 2 | 2 |
| P2 | 6 | 1 | 3 |
| P3 | 3 | 1 | 4 |
| P4 | 4 | 2 | 2 |

Claim matrix C

| | R1 | R2 | R3 |
|----|----|----|----|
| P1 | 1 | 0 | 0 |
| P2 | 6 | 1 | 2 |
| P3 | 2 | 1 | 1 |
| P4 | 0 | 0 | 2 |

Allocation matrix A

| | R1 | R2 | R3 |
|----|----|----|----|
| P1 | 2 | 2 | 2 |
| P2 | 0 | 0 | 1 |
| P3 | 1 | 0 | 3 |
| P4 | 4 | 2 | 0 |

C - A

| R1 | R2 | R3 |
|----|----|----|
| 9 | 3 | 6 |


Resource vector R

| R1 | R2 | R3 |
|----|----|----|
| 0 | 1 | 1 |

Available vector V

(a) Initial state

Determination of a Safe State



| | R1 | R2 | R3 |
|----|----|----|----|
| P1 | 3 | 2 | 2 |
| P2 | 0 | 0 | 0 |
| P3 | 3 | 1 | 4 |
| P4 | 4 | 2 | 2 |

Claim matrix C

| | R1 | R2 | R3 |
|----|----|----|----|
| P1 | 1 | 0 | 0 |
| P2 | 0 | 0 | 0 |
| P3 | 2 | 1 | 1 |
| P4 | 0 | 0 | 2 |

Allocation matrix A

| | R1 | R2 | R3 |
|----|----|----|----|
| P1 | 2 | 2 | 2 |
| P2 | 0 | 0 | 0 |
| P3 | 1 | 0 | 3 |
| P4 | 4 | 2 | 0 |

C - A

| R1 | R2 | R3 |
|----|----|----|
| 9 | 3 | 6 |


Resource vector R

| R1 | R2 | R3 |
|----|----|----|
| 6 | 2 | 3 |

Available vector V

(b) P2 runs to completion

Determination of a Safe State



| | R1 | R2 | R3 |
|----|----|----|----|
| P1 | 0 | 0 | 0 |
| P2 | 0 | 0 | 0 |
| P3 | 3 | 1 | 4 |
| P4 | 4 | 2 | 2 |

Claim matrix **C**

| | R1 | R2 | R3 |
|----|----|----|----|
| P1 | 0 | 0 | 0 |
| P2 | 0 | 0 | 0 |
| P3 | 2 | 1 | 1 |
| P4 | 0 | 0 | 2 |

Allocation matrix **A**

| | R1 | R2 | R3 |
|----|----|----|----|
| P1 | 0 | 0 | 0 |
| P2 | 0 | 0 | 0 |
| P3 | 1 | 0 | 3 |
| P4 | 4 | 2 | 0 |

C - A

| R1 | R2 | R3 |
|----|----|----|
| 9 | 3 | 6 |

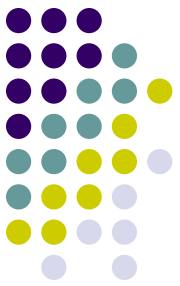
Resource vector **R**

| R1 | R2 | R3 |
|----|----|----|
| 7 | 2 | 3 |

Available vector **V**

(c) **P1 runs to completion**

Determination of a Safe State



| | R1 | R2 | R3 |
|----|----|----|----|
| P1 | 0 | 0 | 0 |
| P2 | 0 | 0 | 0 |
| P3 | 0 | 0 | 0 |
| P4 | 4 | 2 | 2 |

Claim matrix **C**

| | R1 | R2 | R3 |
|----|----|----|----|
| P1 | 0 | 0 | 0 |
| P2 | 0 | 0 | 0 |
| P3 | 0 | 0 | 0 |
| P4 | 0 | 0 | 2 |

Allocation matrix **A**

| | R1 | R2 | R3 |
|----|----|----|----|
| P1 | 0 | 0 | 0 |
| P2 | 0 | 0 | 0 |
| P3 | 0 | 0 | 0 |
| P4 | 4 | 2 | 0 |

C - A

| R1 | R2 | R3 |
|----|----|----|
| 9 | 3 | 6 |

Resource vector **R**

| R1 | R2 | R3 |
|----|----|----|
| 9 | 3 | 4 |

Available vector **V**

(d) P3 runs to completion

Determination of an Unsafe State



| | R1 | R2 | R3 |
|----|----|----|----|
| P1 | 3 | 2 | 2 |
| P2 | 6 | 1 | 3 |
| P3 | 3 | 1 | 4 |
| P4 | 4 | 2 | 2 |

Claim matrix **C**

| | R1 | R2 | R3 |
|----|----|----|----|
| P1 | 1 | 0 | 0 |
| P2 | 5 | 1 | 1 |
| P3 | 2 | 1 | 1 |
| P4 | 0 | 0 | 2 |

Allocation matrix **A**

| | R1 | R2 | R3 |
|----|----|----|----|
| P1 | 2 | 2 | 2 |
| P2 | 1 | 0 | 2 |
| P3 | 1 | 0 | 3 |
| P4 | 4 | 2 | 0 |

C - A

| R1 | R2 | R3 |
|----|----|----|
| 9 | 3 | 6 |

Resource vector **R**

| R1 | R2 | R3 |
|----|----|----|
| 1 | 1 | 2 |

Available vector **V**

(a) Initial state

| | R1 | R2 | R3 |
|----|----|----|----|
| P1 | 3 | 2 | 2 |
| P2 | 6 | 1 | 3 |
| P3 | 3 | 1 | 4 |
| P4 | 4 | 2 | 2 |

Claim matrix **C**

| | R1 | R2 | R3 |
|----|----|----|----|
| P1 | 2 | 0 | 1 |
| P2 | 5 | 1 | 1 |
| P3 | 2 | 1 | 1 |
| P4 | 0 | 0 | 2 |

Allocation matrix **A**

| | R1 | R2 | R3 |
|----|----|----|----|
| P1 | 1 | 2 | 1 |
| P2 | 1 | 0 | 2 |
| P3 | 1 | 0 | 3 |
| P4 | 4 | 2 | 0 |

C - A

| R1 | R2 | R3 |
|----|----|----|
| 9 | 3 | 6 |

Resource vector **R**

| R1 | R2 | R3 |
|----|----|----|
| 0 | 1 | 1 |

Available vector **V**

(b) P1 requests one unit each of R1 and R3

Deadlock Avoidance Logic



```
struct state {  
    int resource[m];  
    int available[m];  
    int claim[n][m];  
    int alloc[n][m];  
}
```

(a) global data structures

```
if (alloc [i,*] + request [*] > claim [i,*])  
    < error >;                                /* total request > claim*/  
else if (request [*] > available [*])  
    < suspend process >;  
else {                                          /* simulate alloc */  
    < define newstate by:  
        alloc [i,*] = alloc [i,*] + request [*];  
        available [*] = available [*] - request [*] >;  
    }  
    if (safe (newstate))  
        < carry out allocation >;  
    else {  
        < restore original state >;  
        < suspend process >;  
    }
```

(b) resource alloc algorithm

Deadlock Avoidance Logic



```
boolean safe (state S) {
    int currentavail[m];
    process rest[<number of processes>];
    currentavail = available;
    rest = {all processes};
    possible = true;
    while (possible) {
        <find a process  $P_k$  in rest such that
            claim  $[k,*] - \text{alloc } [k,*] \leq \text{currentavail};>$ 
        if (found) {                                /* simulate execution of  $P_k$  */
            currentavail = currentavail + alloc  $[k,*]$ ;
            rest = rest -  $\{P_k\}$ ;
        }
        else possible = false;
    }
    return (rest == null);
}
```

(c) test for safety algorithm (banker's algorithm)

Figure 6.9 Deadlock Avoidance Logic

Deadlock Avoidance



- Maximum resource requirement must be stated in advance
- Processes under consideration must be independent; no synchronization requirements
- There must be a fixed number of resources to allocate
- No process may exit while holding resources

Deadlock Detection



| | R1 | R2 | R3 | R4 | R5 |
|----|----|----|----|----|----|
| P1 | 0 | 1 | 0 | 0 | 1 |
| P2 | 0 | 0 | 1 | 0 | 1 |
| P3 | 0 | 0 | 0 | 0 | 1 |
| P4 | 1 | 0 | 1 | 0 | 1 |

Request matrix Q

| | R1 | R2 | R3 | R4 | R5 |
|----|----|----|----|----|----|
| P1 | 1 | 0 | 1 | 1 | 0 |
| P2 | 1 | 1 | 0 | 0 | 0 |
| P3 | 0 | 0 | 0 | 1 | 0 |
| P4 | 0 | 0 | 0 | 0 | 0 |

Allocation matrix A

| R1 | R2 | R3 | R4 | R5 |
|----|----|----|----|----|
| 2 | 1 | 1 | 2 | 1 |

Resource vector

| R1 | R2 | R3 | R4 | R5 |
|----|----|----|----|----|
| 0 | 0 | 0 | 0 | 1 |

Allocation vector

Figure 6.10 Example for Deadlock Detection

Strategies Once Deadlock Detected



- . Abort all deadlocked processes
- . Back up each deadlocked process to some previously defined checkpoint, and restart all process
 - . Original deadlock may occur

Strategies Once Deadlock Detected



- Successively abort deadlocked processes until deadlock no longer exists
- Successively preempt resources until deadlock no longer exists

Advantages and Disadvantages



Table 6.1 Summary of Deadlock Detection, Prevention, and Avoidance Approaches for Operating Systems [ISLO80]

| Approach | Resource Allocation Policy | Different Schemes | Major Advantages | Major Disadvantages |
|------------|--|---|---|---|
| Prevention | Conservative; undercommits resources | Requesting all resources at once | <ul style="list-style-type: none"> • Works well for processes that perform a single burst of activity • No preemption necessary | <ul style="list-style-type: none"> • Inefficient • Delays process initiation • Future resource requirements must be known by processes |
| | | Preemption | <ul style="list-style-type: none"> • Convenient when applied to resources whose state can be saved and restored easily | <ul style="list-style-type: none"> • Preempts more often than necessary |
| | | Resource ordering | <ul style="list-style-type: none"> • Feasible to enforce via compile-time checks • Needs no run-time computation since problem is solved in system design | <ul style="list-style-type: none"> • Disallows incremental resource requests |
| Avoidance | Midway between that of detection and prevention | Manipulate to find at least one safe path | <ul style="list-style-type: none"> • No preemption necessary | <ul style="list-style-type: none"> • Future resource requirements must be known by OS • Processes can be blocked for long periods |
| Detection | Very liberal; requested resources are granted where possible | Invoke periodically to test for deadlock | <ul style="list-style-type: none"> • Never delays process initiation • Facilitates online handling | <ul style="list-style-type: none"> • Inherent preemption losses |

Dining Philosophers Problem





Figure 6.11 Dining Arrangement for Philosophers

Dining Philosophers Problem



```
/* program      diningphilosophers */
semaphore fork [5] = {1};
int i;
void philosopher (int i)
{
    while (true) {
        think();
        wait (fork[i]);
        wait (fork [(i+1) mod 5]);
        eat();
        signal(fork [(i+1) mod 5]);
        signal(fork[i]);
    }
}
void main()
{
    parbegin (philosopher (0), philosopher (1), philosopher
(2),
            philosopher (3), philosopher (4));
}
```

Figure 6.12 A First Solution to the Dining Philosophers Problem

Dining Philosophers Problem



```
/* program diningphilosophers */
semaphore fork[5] = {1};
semaphore room = {4};
int i;
void philosopher (int i)
{
    while (true) {
        think();
        wait (room);
        wait (fork[i]);
        wait (fork [(i+1) mod 5]);
        eat();
        signal (fork [(i+1) mod 5]);
        signal (fork[i]);
        signal (room);
    }
}
void main()
{
    parbegin (philosopher (0), philosopher (1), philosopher (2),
              philosopher (3), philosopher (4));
}
```

Figure 6.13 A Second Solution to the Dining Philosophers Problem

Dining Philosophers Problem



```
monitor dining_controller;
cond ForkReady[5];          /* condition variable for synchronization */
boolean fork[5] = {true};   /* availability status of each fork */

void get_forks(int pid)      /* pid is the philosopher id number */
{
    int left = pid;
    int right = (++pid) % 5;
    /*grant the left fork*/
    if (!fork(left))
        cwait(ForkReady[left]);          /* queue on condition variable */
    fork(left) = false;
    /*grant the right fork*/
    if (!fork(right))
        cwait(ForkReady[right]);         /* queue on condition variable */
    fork(right) = false;
}

void release_forks(int pid)
{
    int left = pid;
    int right = (++pid) % 5;
    /*release the left fork*/
    if (empty(ForkReady[left])          /*no one is waiting for this fork */
        fork(left) = true;
    else                                /* awaken a process waiting on this fork */
        csignal(ForkReady[left]);
    /*release the right fork*/
    if (empty(ForkReady[right])         /*no one is waiting for this fork */
        fork(right) = true;
    else                                /* awaken a process waiting on this fork */
        csignal(ForkReady[right]);
}
```

Dining Philosophers Problem



```
void philosopher[k=0 to 4]          /* the five philosopher clients */
{
    while (true) {
        <think>;
        get forks(k);                /* client requests two forks via monitor */
        <eat spaghetti>;
        release forks(k);            /* client releases forks via the monitor */
    }
}
```

Figure 6.14 A Solution to the Dining Philosophers Problem Using a Monitor

UNIX Concurrency Mechanisms



- Pipes
- Messages
- Shared memory
- Semaphores
- Signals

UNIX Signals



| Value | Name | Description |
|-------|---------|--|
| 01 | SIGHUP | Hang up; sent to process when kernel assumes that the user of that process is doing no useful work |
| 02 | SIGINT | Interrupt |
| 03 | SIGQUIT | Quit; sent by user to induce halting of process and production of core dump |
| 04 | SIGILL | Illegal instruction |
| 05 | SIGTRAP | Trace trap; triggers the execution of code for process tracing |
| 06 | SIGIOT | IOT instruction |
| 07 | SIGEMT | EMT instruction |
| 08 | SIGFPE | Floating-point exception |
| 09 | SIGKILL | Kill; terminate process |
| 10 | SIGBUS | Bus error |
| 11 | SIGSEGV | Segmentation violation; process attempts to access location outside its virtual address space |
| 12 | SIGSYS | Bad argument to system call |
| 13 | SIGPIPE | Write on a pipe that has no readers attached to it |
| 14 | SIGALRM | Alarm clock; issued when a process wishes to receive a signal after a period of time |
| 15 | SIGTERM | Software termination |
| 16 | SIGUSR1 | User-defined signal 1 |
| 17 | SIGUSR2 | User-defined signal 2 |
| 18 | SIGCHLD | Death of a child |
| 19 | SIGPWR | Power failure |

Linux Kernel Concurrency Mechanism



- Includes all the mechanisms found in UNIX
- Atomic operations execute without interruption and without interference


Linux Atomic Operations



Table 6.3 Linux Atomic Operations


| Atomic Integer Operations | |
|--|--|
| <code>ATOMIC_INIT (int i)</code> | At declaration: initialize an <code>atomic_t</code> to <code>i</code> |
| <code>int atomic_read(atomic_t *v)</code> | Read integer value of <code>v</code> |
| <code>void atomic_set(atomic_t *v, int i)</code> | Set the value of <code>v</code> to integer <code>i</code> |
| <code>void atomic_add(int i, atomic_t *v)</code> | Add <code>i</code> to <code>v</code> |
| <code>void atomic_sub(int i, atomic_t *v)</code> | Subtract <code>i</code> from <code>v</code> |
| <code>void atomic_inc(atomic_t *v)</code> | Add 1 to <code>v</code> |
| <code>void atomic_dec(atomic_t *v)</code> | Subtract 1 from <code>v</code> |
| <code>int atomic_sub_and_test(int i, atomic_t *v)</code> | Subtract <code>i</code> from <code>v</code> ; return 1 if the result is zero; return 0 otherwise |
| <code>int atomic_add_negative(int i, atomic_t *v)</code> | Add <code>i</code> to <code>v</code> ; return 1 if the result is negative; return 0 otherwise (used for implementing semaphores) |
| <code>int atomic_dec_and_test(atomic_t *v)</code> | Subtract 1 from <code>v</code> ; return 1 if the result is zero; return 0 otherwise |
| <code>int atomic_inc_and_test(atomic_t *v)</code> | Add 1 to <code>v</code> ; return 1 if the result is zero; return 0 otherwise |

Linux Atomic Operations



| Atomic Bitmap Operations | |
|--|---|
| <code>void set_bit(int nr, void *addr)</code> | Set bit nr in the bitmap pointed to by addr |
| <code>void clear_bit(int nr, void *addr)</code> | Clear bit nr in the bitmap pointed to by addr |
| <code>void change_bit(int nr, void *addr)</code> | Invert bit nr in the bitmap pointed to by addr |
| <code>int test_and_set_bit(int nr, void *addr)</code> | Set bit nr in the bitmap pointed to by addr; return the old bit value |
| <code>int test_and_clear_bit(int nr, void *addr)</code> | Clear bit nr in the bitmap pointed to by addr; return the old bit value |
| <code>int test_and_change_bit(int nr, void *addr)</code> | Invert bit nr in the bitmap pointed to by addr; return the old bit value |
| <code>int test_bit(int nr, void *addr)</code> | Return the value of bit nr in the bitmap pointed to by addr |

Linux Spinlocks



| | |
|---|---|
| <code>void spin_lock(spinlock_t *lock)</code> | Acquires the specified lock, spinning if needed until it is available |
| <code>void spin_lock_irq(spinlock_t *lock)</code> | Like <code>spin_lock</code> , but also disables interrupts on the local processor |
| <code>void spin_lock_irqsave(spinlock_t *lock, unsigned long flags)</code> | Like <code>spin_lock_irq</code> , but also saves the current interrupt state in flags |
| <code>void spin_lock_bh(spinlock_t *lock)</code> | Like <code>spin_lock</code> , but also disables the execution of all bottom halves |
| <code>void spin_unlock(spinlock_t *lock)</code> | Releases given lock |
| <code>void spin_unlock_irq(spinlock_t *lock)</code> | Releases given lock and enables local interrupts |
| <code>void spin_unlock_irqrestore(spinlock_t *lock, unsigned long flags)</code> | Releases given lock and restores local interrupts to given previous state |
| <code>void spin_unlock_bh(spinlock_t *lock)</code> | Releases given lock and enables bottom halves |
| <code>void spin_lock_init(spinlock_t *lock)</code> | Initializes given spinlock |
| <code>int spin_trylock(spinlock_t *lock)</code> | Tries to acquire specified lock; returns nonzero if lock is currently held and zero otherwise |
| <code>int spin_is_locked(spinlock_t *lock)</code> | Returns nonzero if lock is currently held and zero otherwise |

Linux Semaphores



| Traditional Semaphores | |
|---|---|
| <code>void sema_init(struct semaphore *sem, int count)</code> | Initializes the dynamically created semaphore to the given count |
| <code>void init_MUTEX(struct semaphore *sem)</code> | Initializes the dynamically created semaphore with a count of 1 (initially unlocked) |
| <code>void init_MUTEX_LOCKED(struct semaphore *sem)</code> | Initializes the dynamically created semaphore with a count of 0 (initially locked) |
| <code>void down(struct semaphore *sem)</code> | Attempts to acquire the given semaphore, entering uninterruptible sleep if semaphore is unavailable |
| <code>int down_interruptible(struct semaphore *sem)</code> | Attempts to acquire the given semaphore, entering interruptible sleep if semaphore is unavailable; returns -EINTR value if a signal other than the result of an up operation is received. |
| <code>int down_trylock(struct semaphore *sem)</code> | Attempts to acquire the given semaphore, and returns a nonzero value if semaphore is unavailable |
| <code>void up(struct semaphore *sem)</code> | Releases the given semaphore |
| Reader-Writer Semaphores | |
| <code>void init_rwsem(struct rw_semaphore, *rwsem)</code> | Initializes the dynamically created semaphore with a count of 1 |
| <code>void down_read(struct rw_semaphore, *rwsem)</code> | Down operation for readers |
| <code>void up_read(struct rw_semaphore, *rwsem)</code> | Up operation for readers |
| <code>void down_write(struct rw_semaphore, *rwsem)</code> | Down operation for writers |
| <code>void up_write(struct rw_semaphore, *rwsem)</code> | Up operation for writers |

Linux Memory Barrier Operations



Table 6.6 Linux Memory Barrier Operations

| | |
|--------------------------|---|
| <code>rmb ()</code> | Prevents loads from being reordered across the barrier |
| <code>wmb ()</code> | Prevents stores from being reordered across the barrier |
| <code>mb ()</code> | Prevents loads and stores from being reordered across the barrier |
| <code>Barrier ()</code> | Prevents the compiler from reordering loads or stores across the barrier |
| <code>smp_rmb ()</code> | On SMP, provides a <code>rmb ()</code> and on UP provides a <code>barrier ()</code> |
| <code>smp_wmb ()</code> | On SMP, provides a <code>wmb ()</code> and on UP provides a <code>barrier ()</code> |
| <code>smp_mb ()</code> | On SMP, provides a <code>mb ()</code> and on UP provides a <code>barrier ()</code> |

SMP = symmetric multiprocessor

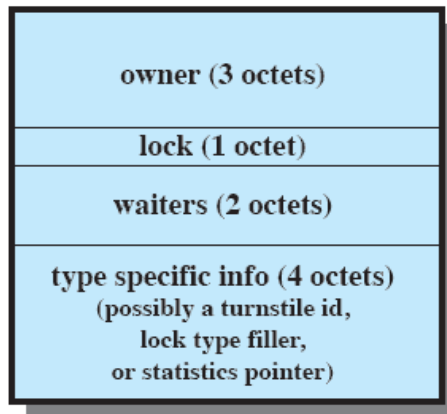
UP = uniprocessor

Solaris Thread Synchronization Primitives




- Mutual exclusion (mutex) locks
- Semaphores
- Multiple readers, single writer (readers/writer) locks
- Condition variables

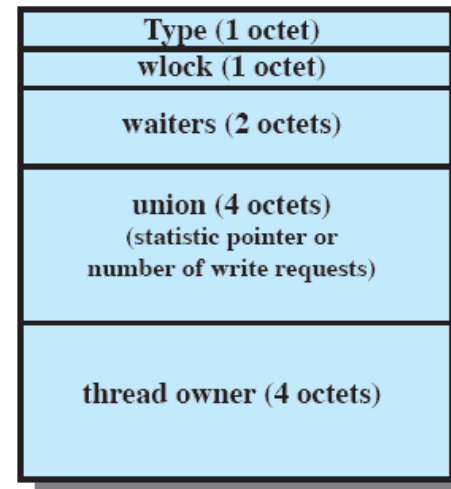
Solaris Synchronization Data Structures



(a) MUTEX lock



(b) Semaphore



(c) Reader/writer lock



(d) Condition variable

Figure 6.15 Solaris Synchronization Data Structures

Windows Synchronization Objects



| Object Type | Definition | Set to Signaled State When | Effect on Waiting Threads |
|-----------------------|---|--|---------------------------|
| Notification Event | An announcement that a system event has occurred | Thread sets the event | All released |
| Synchronization event | An announcement that a system event has occurred. | Thread sets the event | One thread released |
| Mutex | A mechanism that provides mutual exclusion capabilities; equivalent to a binary semaphore | Owning thread or other thread releases the mutex | One thread released |
| Semaphore | A counter that regulates the number of threads that can use a resource | Semaphore count drops to zero | All released |
| Waitable timer | A counter that records the passage of time | Set time arrives or time interval expires | All released |
| File | An instance of an opened file or I/O device | I/O operation completes | All released |
| Process | A program invocation, including the address space and resources required to run the program | Last thread terminates | All released |
| Thread | An executable entity within a process | Thread terminates | All released |

Note: Shaded rows correspond to objects that exist for the sole purpose of synchronization.